GCVE-BCP-02 - Practical Guide to Vulnerability Handling and Disclosure

GCVE.eu

GCVE.eu

Contents

1 Practical Guide to Vulnerability Handling and Disclosure **1**

1.1 Introduction . . . 1

1.1.1 Vulnerability Handling at a Glance . . . 2

1.2 Definitions . . . 3

1.3 Roles and Responsibilities . . . 3

1.4 Preparing for Vulnerability Handling . . . 5

1.4.1 Preparation Overview . . . 7

1.5 Receiving and Handling Vulnerability Reports . . . 8

1.5.1 Finding Issues Before a Report Arrives . . . 10

1.5.2 Report Intake and Verification Overview . . . 10

1.6 Investigating and Resolving Vulnerabilities . . . 11

1.6.1 Getting Security Updates to Users . . . 13

1.6.2 Investigation, Remediation, and Delivery Overview . . . 13

1.7 Communicating with Submitters and Users . . . 14

1.7.1 Communication Overview . . . 17

1.8 Coordinated Vulnerability Disclosure . . . 17

1.8.1 Coordinated Disclosure Decision Overview . . . 19

1.9 Publishing Advisories . . . 21

1.9.1 Advisory Publication and Follow-up Overview . . . 24

1.10 Best Practices . . . 25

2 Appendix A: EU Cyber Resilience Act Reporting **27**

2.1 What the CRA Adds . . . 27

2.2 Relationship to Coordinated Vulnerability Disclosure . . . 27

2.3 The Single Reporting Platform . . . 28

2.3.1 Reporting stages and deadlines . . . 29

2.4 Integrating CRA Reporting into a CVD Process . . . 29

2.4.1 Parallel CRA Reporting and CVD Overview . . . 30

3 References **33**

4 Acknowledgements **35**
4.1 BCP-02 Coordinators 35
4.2 Contributions 35

1 Practical Guide to Vulnerability Handling and Disclosure



GCVE.eu

- **Version**: 1.9
- **Status**: Published
- **Date**: 2026-10-15
- **Authors**: GCVE Working Group
- **BCP ID**: BCP-02

This guide is distributed and available under CC-BY-4.0.

Copyright (C) 2025-2026 GCVE Initiative.

1.1 Introduction

Vulnerabilities in software can pose serious risks to users and organizations. A clear and effective process for handling and disclosing security vulnerabilities maintains user trust and protects systems.

This guide provides actionable recommendations for GCVE GNA, software developers, open source project maintainers, vendors, and organizations to manage vulnerability reports from discovery to resolution and public disclosure. It is organized into key stages of a vulnerability's life-cycle, including:

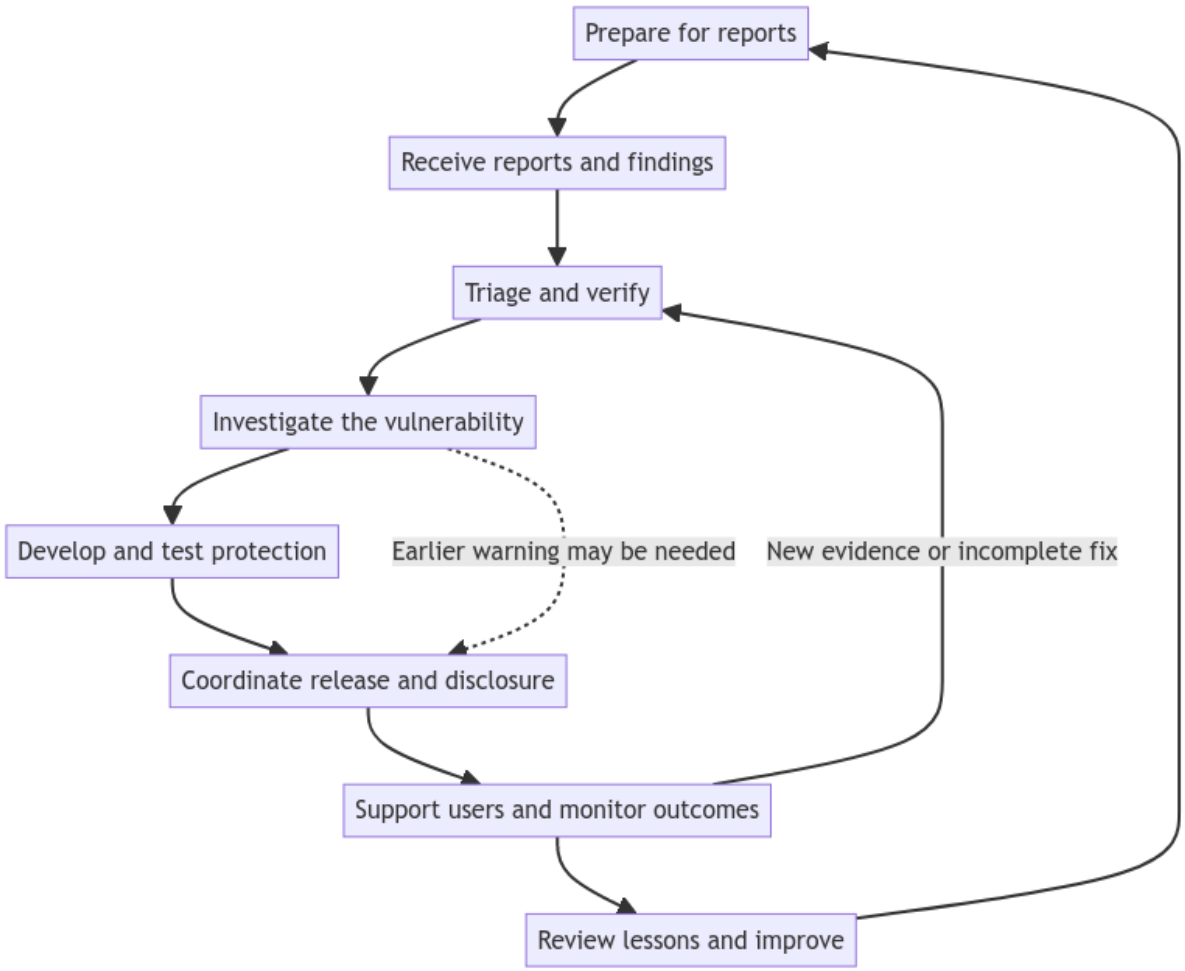
- **preparation and receipt of a report,**

- **investigation and remediation,**
- **communication and coordinated disclosure.**

Overall, the guide establishes a transparent process that encourages responsible reporting and safeguards users.

1.1.1 Vulnerability Handling at a Glance

The following overview connects the stages described in this guide. Communication and coordination continue throughout the case; investigation, remediation, and disclosure may overlap. Where applicable, CRA reporting runs in parallel (see Appendix A).



1.2 Definitions

Vulnerability A flaw or weakness in a software system that could be exploited to compromise the system's security or functionality.

Vulnerability Report Information provided about a potential vulnerability in a product or service. A report typically includes details needed to reproduce the issue, such as affected components, steps to trigger the bug, and expected versus actual behaviour.

Coordinated Vulnerability Disclosure (CVD) A practice where the reporter, vendor, and other affected parties work together to investigate a vulnerability, prepare fixes or mitigations, and agree on public disclosure. Details are usually shared privately while a fix is prepared. Disclosure may still be needed before a fix is available, for example when exploitation is already happening or users need to take protective action.

Advisory A public notice describing a vulnerability, its impact, affected versions, and available fixes or mitigations. If no fix is available, the advisory should say so and explain what users can do to reduce their exposure.

1.3 Roles and Responsibilities

Effective vulnerability handling involves several stakeholders with distinct roles. Defining responsibilities for each role helps ensure accountability and clarity during the process:

- **Submitter or Reporter (Security Researchers or Users):** The person or team who discovers a potential vulnerability and reports it to the appropriate party (usually the vendor or maintainer). This could be an independent security researcher, a user, a member of the development team, or a third-party analyst. Submitters or reporters are expected to investigate and disclose vulnerabilities in good faith.
- **Submitter/Reporter and Vendor Interaction:** The submitter or reporter should follow the vendor's reporting guidelines, avoid publicizing the issue before giving the vendor a chance to fix it, and provide sufficient detail to reproduce the problem. Ethical reporters respect privacy and legal boundaries during testing, and they do not exploit the vulnerability for personal gain. When participating in a bug bounty program, they adhere to its scope and rules.
- **Vendors / Maintainers**: Vendors and open source project maintainers act on vulnerability reports promptly and professionally. They have a clear intake mechanism (e.g. a security contact email or form), acknowledge and maintain open communication with reporters, prioritize, analyse and fix issues based on risk, and communicate updates about confirmed vulnerabilities and their fixes to all users. Vendors must also avoid hostile responses like threatening legal action against those who report issues in good faith.

- **Product Security Team (PSIRT)**: An organization's PSIRT or designated security response team, when present, coordinates the entire handling of the vulnerability response process. They receive and triage incoming reports, involve relevant development teams, track remediation progress, and ensure communication flows to reporters and management. The PSIRT also works on preparing advisories and post-remediation follow-ups. In an open source context, the PSIRT may simply consist of the core maintainers or a subset of project contributors who handle security issues.

- **Developers and QA Engineers**: Once a report is validated, developers are responsible for developing the code fixes or mitigations. QA or testing teams (or in smaller projects, the developers themselves) must test the fixes to confirm the vulnerability is resolved and that no new issues are introduced. They also verify that the fix doesn't negatively impact functionality. Development and QA staff should treat security fixes with high priority and follow secure coding practices to prevent re-introduction of similar flaws.

- **Secure Coding Practices**: Secure coding is the practice of developing computer software in such a way that guards against the accidental introduction of security vulnerabilities. Defects, bugs and logic flaws are consistently the primary cause of commonly exploited software vulnerabilities.

- **Coordinators (Third-Party Facilitators)**: In some cases, a coordinator such as a Computer Emergency Response team (CERT) or a vulnerability broker can ensure responsible information sharing and awareness among multiple vendors impacted by a vulnerability. They facilitate information sharing when the same issue affects multiple vendors or when a reporter and vendor require a neutral intermediary. Coordinators help synchronize remediation timelines among all affected parties and can assist in broader communications, such as publishing a joint advisory.

- **Users/Customers**: While users are the beneficiaries of the process rather than active participants in handling, they are ultimately responsible for applying updates and mitigations that vendors release. Vendors should simplify the process for users to learn about security updates and take action. Users, on their part, should stay informed via the channels the vendor provides (mailing lists, security pages, etc.) and promptly install security patches to protect their systems.

By clearly identifying these roles and expectations, an organization or project can ensure well-managed contributions towards swiftly resolving vulnerabilities when they arise.

1.4 Preparing for Vulnerability Handling

Preparation is critical **before** any vulnerability reports come in. Having a plan and resources in place will make the handling process far more efficient and effective. Key preparation steps include:

- **Establish a Vulnerability Disclosure Policy**: Create a clearly written document that outlines how people should report vulnerabilities to you and what they can expect in return. This policy should be easy to find and may often be a publicly accessible policy on a website or a `SECURITY.md` file in open source projects. It should specify the preferred contact method and provide links to a dedicated security email address or web form. The policy should detail what information to include in a report and outline your commitment to researchers. These commitments should capture any expectations for communication, acknowledgement, and not pursuing legal action for good-faith research. If you have a bug bounty program or reward offering, mention how it works and the scope of vulnerabilities covered. Also include any safe harbour statement reassuring researchers that if they follow the guidelines, they will not be penalized.
- **Designate a Security Response Team**: Assign specific individuals from your organization or project team to the defined roles and responsibilities for handling vulnerability reports. These may be members of the security or engineering team in a company or maintainers in an open source project. These roles should provide for receiving, assessing, and prioritizing reports, fixing the vulnerability, and communicating with stakeholders including the reporter, vendors, and users. Make sure these individuals know their roles and have the authority to coordinate fixes across development teams.

  Use a shared contact address and name a backup so reports do not wait for one person to return. Agree who can approve an urgent release or a disclosure decision. A small project can combine these roles; the important point is that someone owns the next action.
- **Define an Internal Workflow**: Establish what to do when a report arrives:
  - Communicate issue Routing: Everyone in the organization should know how to escalate a security issue to the response team. For example, even customer support or developers who accidentally receive a report should know where to forward it.
  - Set Up Secure Communication Channels: Provide secure and easy-to-use reporting channels. Commonly, a security contact email (e.g. security@example.com) should be combined with a public PGP/GPG key so that sensitive vulnerability details can be encrypted in-transit. Web-based submission forms or ticket systems that restrict reports from public view and use encryption (HTTPS) can also provide a secure channel. If using an issue tracker like those in GitHub or GitLab, instruct reporters to not post vulnerabilities in public tickets; instead provide a private reporting mechanism or use the platform's private disclosure feature if available.

  - Develop Internal Procedures and Guidelines: Document how your team will handle vulnerability reports step by step. This internal guide should include at a minimum the following items:
  - Document triage criteria. Provide an internal guide for how to prioritize reports. Severity rating systems like CVSS can help categorize issues.
  - Publish target timelines. Each step in handling process should have a commitment for resolution, such as acknowledging report receipt within two business days and providing a plan for resolution within one week. Critically, a commitment should be posted about resolving vulnerabilities within a specific time-frame.
  - Establish the technical process: Identify how you will analyse vulnerability reports, develop and test fixes, and deploy patches.
  - Maintain accountability: Set up a controlled and auditable system to track vulnerability cases from report to resolution – this could be a special project in your issue tracker or a simple internal spreadsheet/ticket system. Tracking helps ensure nothing falls through the cracks and allows management to review metrics like response times.

- **Know What You Ship**: Keep a component list for each supported release, including third-party libraries, bundled or copied code, and dependencies pulled in by other dependencies where known. Record names, versions, maintainers or suppliers, and useful identifiers such as package URLs. A machine-readable software bill of materials (SBOM), for example in SPDX or CycloneDX, makes comparison with vulnerability feeds easier. Keep it with the release, update it for new builds, and correct errors when found. Mark gaps in the inventory rather than assuming it is complete. For devices, also record hardware models, revisions, and firmware versions. Consider build tools that could affect the shipped product.
- **Keep Contacts and Support Information Current**: Know how to reach the maintainers of components you use and the projects, distributors, or integrators that depend on your product. Record which releases you support and where users can find that information. Watch for dependencies losing security support or becoming unmaintained, and plan their replacement before an urgent fix depends on them.
- **Plan Regular Security Checks**: Do not rely only on external reports. Set a testing and review schedule that fits the product's exposure and your resources. Combine dependency checks with activities such as code review, fuzzing, or penetration testing where useful. Revisit the plan when the product, its deployment, or known attack techniques change. Feed findings into the same handling process as external reports.
- **Prepare Update Delivery**: Decide how users will receive security updates and advisories before an emergency. Provide a way to verify that updates come from you and have not been

altered. Test the download, package repository, or automatic update path, and keep a manual path where users need one. Link the advisory and installation instructions from the update channel.

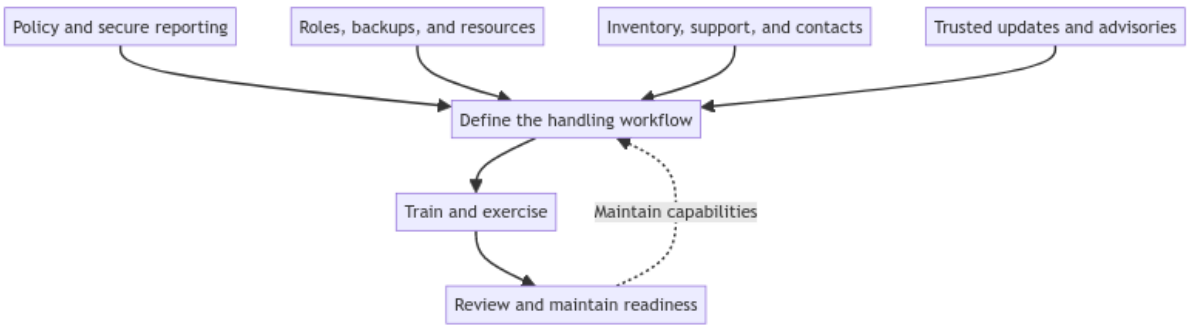
- **Practice and Training**: Treat vulnerability handling as a process that benefits from practice. Run simulations or tabletop exercises that reinforce what workflows to follow when someone reports a critical software bug. Ensure the response team is familiar with the procedures and tools, as many items like encrypted emails and coordinating CVE requests are not frequently used. Train developers on secure coding and common vulnerability types so that they can respond more effectively and ideally prevent issues. Commit to continuous improvement through updating your plan periodically based on these dry runs or any real incidents.

- **Resource Allocation**: Management should allocate adequate resources for vulnerability handling. This means having people with time to investigate and fix security issues on short notice and possibly budgeting for external help if needed. For instance, payments for penetration testing or bug bounties may require financing. Without proper resourcing and clear management support, even a well-documented process can falter. Make sure everyone prioritizes security bug fixes as a part of the development life-cycle rather than as an afterthought.

Advance preparations create a foundation that makes the vulnerability handling workflow run smoothly. It signals to researchers and users alike that your organization takes security seriously and can readily respond.

1.4.1 Preparation Overview

Put these capabilities in place before a report arrives, then exercise and maintain them. Findings from regular security checks enter the same intake process as external reports.



1.5 Receiving and Handling Vulnerability Reports

Vulnerability reports require effective stewardship and communication regarding their resolution to maintain trust and accountability with the research and user communities. Key management practices include:

- **Logging and Tracking**: As part of handling incoming reports, log each report in your internal tracking system. Record details such as the date received, reporter name/contact, affected product, a short description of the issue, its status (e.g. “under investigation”, “fix in progress”), and any relevant deadlines. For instance, if a reporter has indicated they intend to publish after 90 days, note that date. Tracking helps you manage multiple reports at once and ensures accountability. It also creates a record useful for post-incident review and improving your process over time.

  Assign a case owner and record the next action and update date. Keep links to the evidence, affected releases, assessment decisions, fix and test results, correspondence, and advisory. Record changes to the plan and why they were made. Keep non-public details in a restricted workspace, including attachments and logs; avoid copying them into public tickets or build output.

- **Acknowledge Receipt Quickly**: The **intake** and **initial response** sets the tone for the entire handling process. Upon receiving a vulnerability report, assign a tracking ID for future correspondence and following the report throughout verification and resolution. Acknowledge the report as soon as possible – ideally within 24-48 hours. A simple receipt confirmation that indicates a planned review assures the reporter that you value and will act upon their effort. This can prevent a frustrated researcher from going public prematurely.

- **Ensure Confidentiality**: Treat all vulnerability reports as sensitive information. Restrict knowledge of the details to the people who need to know (your security team, relevant developers, management as appropriate). If a report comes in through a public channel by mistake (e.g. a public bug tracker or social media), move the conversation to a private channel promptly and ask the reporter to take down public details if possible. This limits exposure of the vulnerability until a fix is ready and reduces the window of opportunity for malicious actors.

  Accept and route useful reports even if they arrive through the wrong channel, lack encryption, or do not use your preferred template. Offer a safer channel for follow-up instead of rejecting the report. Keep reporting instructions easy to understand and usable with assistive tools, and offer an alternative contact method when needed.

- **Initial Triage (Severity and Scope Assessment)**: Your security response team should quickly triage the report. Determine what product the vulnerability affects and whether it involves a core component or optional module. Identify the potential vulnerability type and its severity. A

defined severity scoring system like CVSS can isolate the severity of specific vulnerabilities separately from their risk (e.g. does it allow remote code execution, information disclosure, denial of service, etc.). Also assess the scope considering the number of users likely affected. During this phase, it's fine to ask the reporter for clarification or additional details if needed to understand the issue. Based on this initial assessment, you can prioritize the issue. For example, critical issues trigger an emergency response whereas low-risk issues go into the normal development queue.

Use the actual deployment to set priority: is the vulnerable code reachable, exposed to untrusted input, or protected by an effective mitigation? Consider impact on users and evidence of exploitation alongside the severity score. Record the priority and the reason for it, including whether emergency handling is needed, and reassess when new information arrives. Lack of a public exploit does not show that a product is unaffected.

- **Verify the Report**: Before immediately developing a fix, reproduce and verify the issue to confirm that it is a genuine vulnerability. This may involve your engineers replicating the steps provided by the reporter or creating a proof-of-concept exploit in a safe test environment. Verification should also document which versions/configurations are affected and what the exact impact is (e.g. can an attacker actually steal data, or is it a minor glitch?). If the report cannot be verified or appears to be incorrect, involve the reporter – let them know if additional evidence is needed or explain why it might not be considered a security issue. In some cases, what the reporter found might be a known issue or expected behaviour; communicate that respectfully if so.

  Distinguish "not affected" from "not yet reproduced". Before closing a case, record what was checked, the evidence and assumptions behind the decision, and what would justify reopening it. Explain the outcome to the reporter and revisit it if new evidence or a different deployment changes the assessment. Urgent protective action need not wait for every investigation step to finish.

- **Duplicate or Out-of-Scope Issues**: If you discover that a reported issue duplicates something already reported or concerns a product you do not ship or maintain, explain the situation to the reporter. For duplicates, thank them and inform them of the issue's previous resolution. If possible, indicate the tracking number or status of the existing issue. For third-party issues, you might need to act as a coordinator and relay the report to the third-party vendor or to an appropriate coordinating centre. Similarly, keep the reporter advised of this step. If the issue is in an older end-of-life product version that you no longer support, let the reporter know the product is not supported. However, if the vulnerability is critical, you might still consider informing users or updating documentation about the risk.

  A flaw in a library you ship still needs an assessment for your product. Contacting the upstream

maintainer does not replace tracking your own affected releases, integrating a fix or mitigation, and informing your users. Link duplicate reports to the existing case without exposing another reporter's private information.

- **Ongoing Communication**: Throughout the handling of the report, maintain communication with the reporter. After the initial acknowledgement, update them when verification is done to confirm results from the verification or to ask for more information to reproduce the issue. Even if you don't have significant progress, a periodic update (e.g. "We are still working on a fix, thank you for your patience") can reassure reporters. Communication provides transparency and communicates your commitment towards resolving vulnerabilities. This collaborative approach keeps the process coordinated and reduces the likelihood of surprise disclosures.

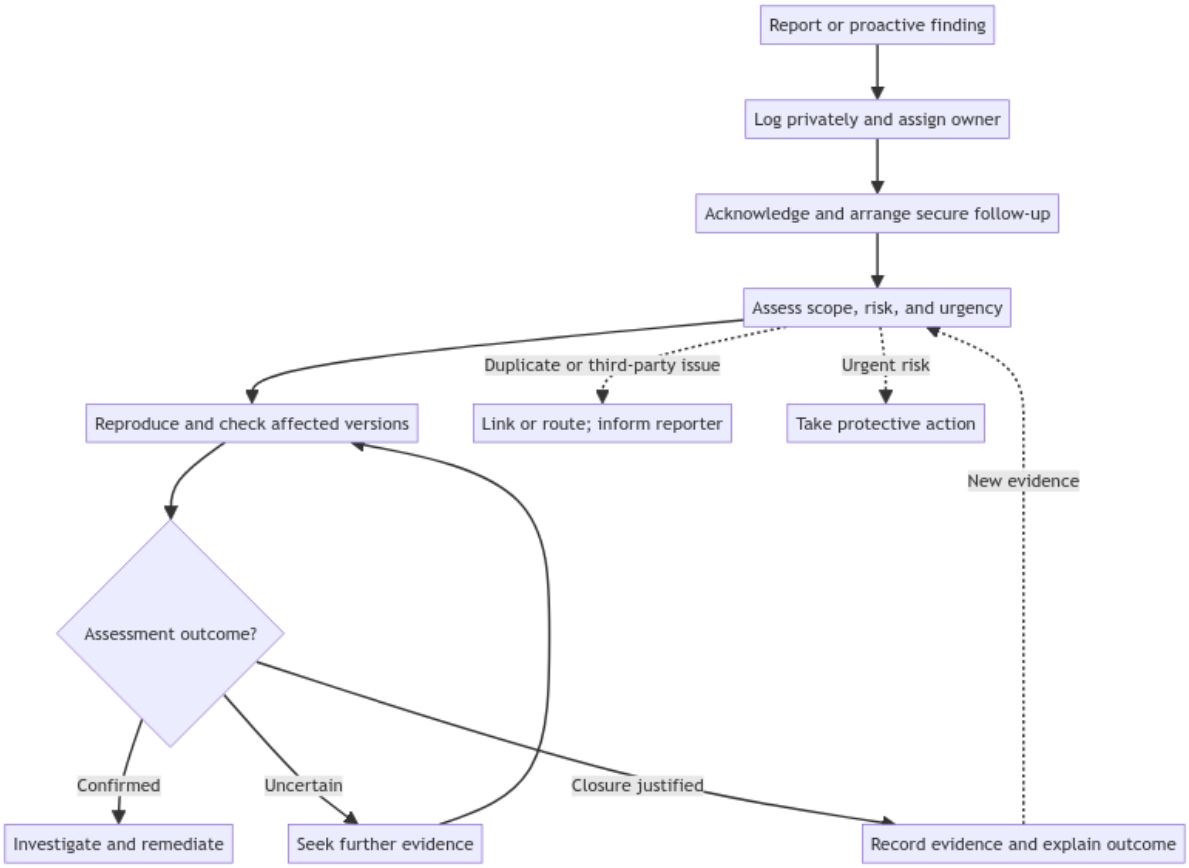
Receiving and handling reports properly builds trust with researchers. If they see you are responsive and professional, they are more likely to continue reporting to you rather than announcing issues publicly out of frustration. It also signals to your user community that you can deal with security issues seriously and efficiently.

1.5.1 Finding Issues Before a Report Arrives

Monitor upstream security advisories and relevant vulnerability sources, such as GCVE feeds and the EU Vulnerability Database (EUVD), continuously or on a schedule suited to the product's risk. Assign someone to review alerts and compare them with the component lists for supported releases. Use automation to find possible matches, then check applicability rather than treating every name or version match as a confirmed vulnerability. Track relevant findings from these sources, internal tests, and security reviews in the same way as incoming reports.

1.5.2 Report Intake and Verification Overview

The response team owns the case and keeps the reporter informed. A duplicate links to the existing case; a component shipped in your product still requires your own impact assessment. A failure to reproduce is a reason to seek evidence, rather than proof that the product is unaffected.



1.6 Investigating and Resolving Vulnerabilities

Once a vulnerability report has been validated, begin investigating its root cause and implementing a resolution. During this phase development and security teams work together to eliminate the weakness. Important steps in investigation and resolution include:

- **In-Depth Analysis**: The team should thoroughly investigate the vulnerability. Root cause evaluations determine which part of the code or design is at fault. Understanding root cause is vital not only to fixing the current bug but also to identifying any related issues. Investigate if similar components might have the same flaw and whether the vulnerability has resided throughout other versions. Also, analyse the impact achieved by an attacker exploiting the vulnerability under different conditions. This helps in assessing severity and priority for the fix . For instance, a vulnerability allowing full system compromise is critical, whereas one causing a minor information leak might be lower priority). Formal severity scoring, like CVSS, can be useful to quantify impact and guide prioritization.

- **Developing a Fix or Mitigation**: Task the appropriate development team to create a remediation for the issue. This often means writing a patch to the software's code. If a full fix might take too much time, consider if a mitigation or workaround could reduce the interim risk. For example, instructing users to disable a vulnerable feature or providing a configuration change that blocks an exploit might be a temporary measure. Balance speed and quality when developing fixes – urgent fixes should be expedited, but even a quick fix needs at least some testing to ensure it resolves the problem without causing regressions. In extreme cases involving an actively exploited critical vulnerability, a vendor might release a temporary fix or even temporarily take a service offline until a proper patch is ready. These emergency measures underscore the importance of having a robust fix as soon as possible.

  Record who will deliver the fix, which supported branches need it, dependencies on other teams, and target dates for testing and release. Keep the security change separate from unrelated features where practical so users can deploy it with less disruption. If separation is not feasible, explain the constraint and any additional upgrade risk.

- **Testing the Fix**: Thorough testing verifies that the fix indeed closes the vulnerability and does not introduce new bugs or break functionality. Testing on all relevant platforms or versions of the software ensures consistency. Where practical, involving the original reporter in testing a patch or a secure test build can provide valuable confirmation that the vulnerability is resolved from an attacker's perspective. This also further engages the reporter in the process. If the reporter is not available for testing or the disclosure is sensitive, rigorous internal testing should leverage initial report descriptions, root causes and conditions identified during verification and development, and edge cases that could be related to the vulnerability.

  Test upstream patches in your own product and configuration; an upstream test result alone does not confirm that your users are protected. Retain the results and add a regression test where practical. Also check installation, compatibility, and significant side effects, including performance or loss of functionality.

- **Collateral Clean-up**: Investigating a vulnerability might reveal other issues or similar vulnerable code elsewhere. Take this opportunity to clean up related problems. For example, if the issue was a result of an insecure library or dependency, update that dependency across your project. If the root cause was a specific coding pattern, a quick scan of the codebase can reveal instances of that pattern. This proactive approach can prevent future vulnerabilities. It's also a good practice to check if any telemetry, incident reports, or logs provide evidence that the vulnerability was exploited. If so, this necessitates incident response.

- **Decision on Public Release Timing**: Plan when and how the fix will be released. Typically, you will coordinate the release of the patched software with the publication of a security advisory (see the sections on Communication and Publishing Advisories). If you have a regular release

schedule, decide if this warrants an out-of-cycle emergency release or if it can wait for the next planned update. Security patches for critical issues often justify quicker releases. If the vulnerability is not very risky, bundling it in a scheduled release might be acceptable. Aim to release the fix before or at the same time as public technical details. If users need an earlier warning, follow the guidance on no-fix scenarios and advisory timing below.

- **Documenting the Fix**: When releasing the fix, prepare internal documentation of what was done. This includes updating any internal security knowledge base about the nature of the issue and how it was fixed for future reference. If a CVE (Common Vulnerabilities and Exposures) or GCVE ID is needed (and not already assigned by a reporter or coordinator), you might request one at this stage. Many open source projects can get CVEs via CNA coordination or GCVE via GNA or through portals like GitHub Security Advisories. One or more ID will be referenced in the public advisory to track the issue in vulnerability databases.

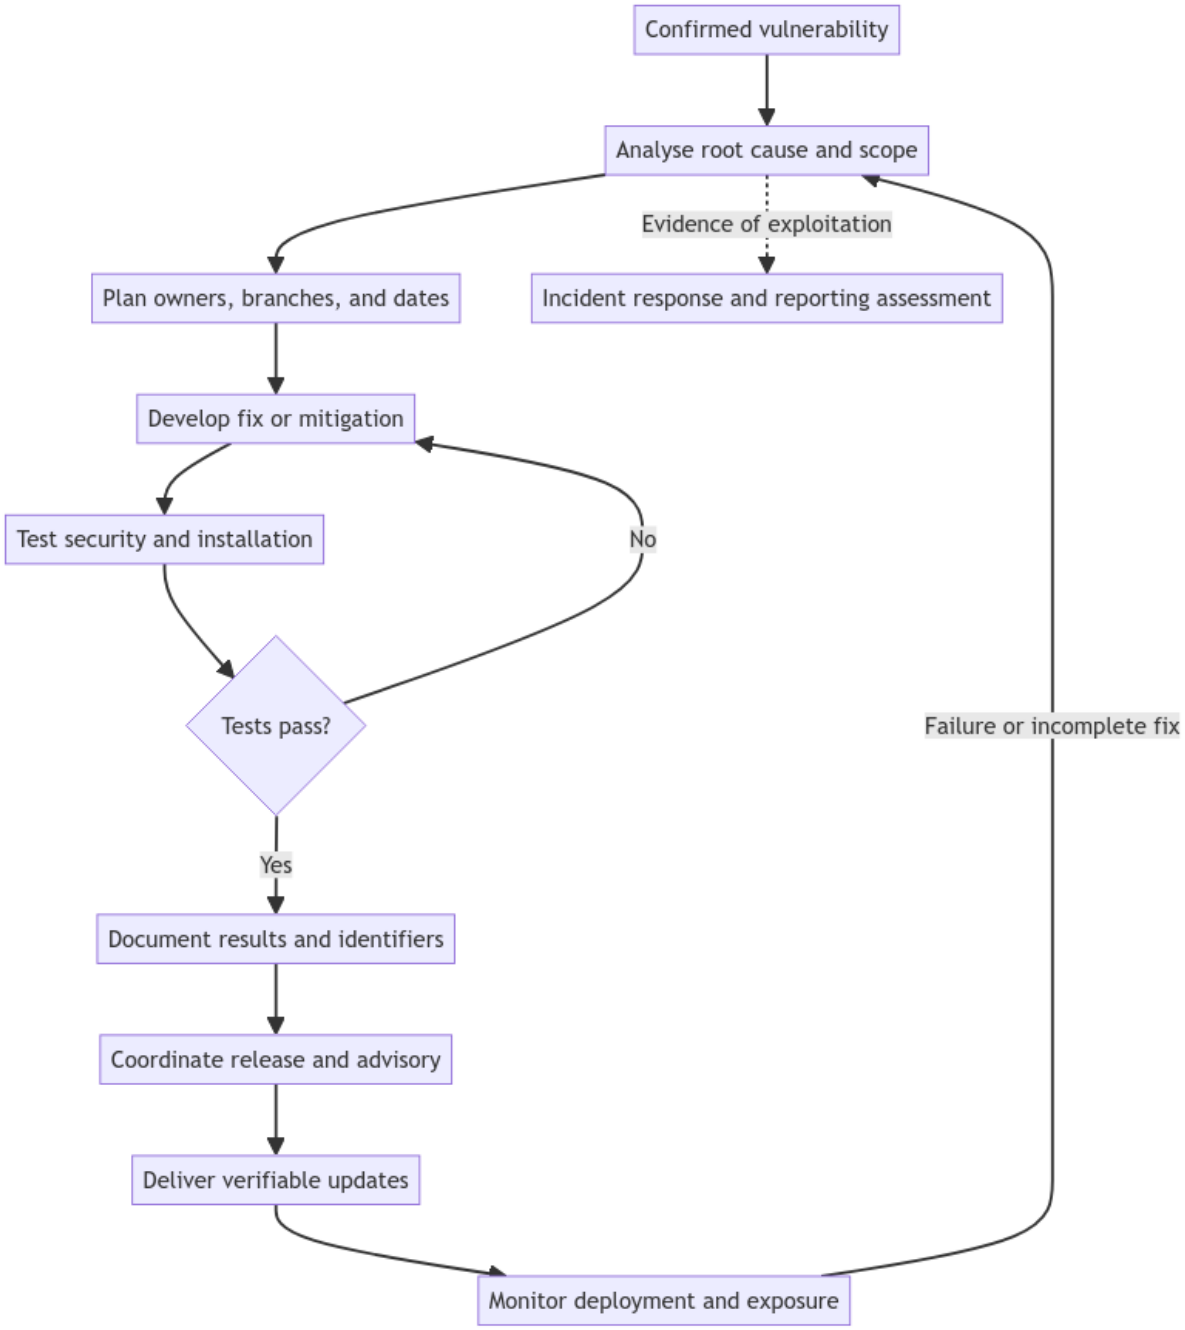
Diligent investigation and efficient remediation eliminate a vulnerability's threat. It's important to keep momentum throughout resolution – lengthy delays in fixing known vulnerabilities leave users at risk. Aim for a balance among scope, severity, and risk: address the issue as fast as possible, but also correctly and safely. Once a fix is ready and verified, begin preparations to communicate both the fix and disclose the vulnerability.

1.6.1 Getting Security Updates to Users

- **Release Through a Trusted Channel**: Use the prepared update path and protect the authenticity and integrity of packages and update metadata, for example with signatures verified by the updater or package manager. Provide clear verification instructions for manual downloads. Offer security updates without an extra charge and avoid making users buy an unrelated upgrade to obtain a security fix.
- **Make Installation Practical**: State prerequisites, restart requirements, configuration changes, and any service interruption. Explain how to check that the fixed version or mitigation is active. For automatic updates, check that delivery and installation succeed; for manual updates, provide reproducible steps. Where an update can fail, provide recovery guidance that makes clear if returning to an older version restores the vulnerability.

1.6.2 Investigation, Remediation, and Delivery Overview

Developers and QA verify both the security correction and its delivery in the affected product. Keep disclosure planning and reporter updates active during this work. If no full fix is available, follow the no-fix and earlier-warning paths in the disclosure section.



1.7 Communicating with Submitters and Users

Communication connects the entire vulnerability handling process. Its management focuses on the vulnerability reporter and the users or customers of the affected product.

- **Communication with the Reporter:** From the moment a report is received until after the issue is resolved, it's best practice to keep the reporter informed and engaged. Key communication milestones with reporters include:
- **Acknowledgement**: As noted earlier, confirm to the reporter that you received their report and provide a reference ID for it. Thank them for alerting you.
- **Verification Results**: After you investigate, let the reporter know the outcome. If you verified the vulnerability, say so and that you've begun remediation efforts. If you could not reproduce the issue or determined it's not a vulnerability, provide that feedback. Be tactful and appreciative – even if a report turns out not to be a valid security issue, acknowledge the effort and explain your reasoning to avoid discouraging future reports.
- **Updates During Remediation**: While working on the fix, send periodic updates. Although these do not need to share full technical details, they should inform reporters of progress towards identifying root causes, developing patches, and testing fixes. If there are delays or unforeseen complications, be honest about them and your continued work towards resolution. Regular communication keeps the reporter on your side and maintains trust.
- **Coordination on Disclosure**: As the fix nears completion, coordinate the public disclosure with the reporter. If you plan to publish an advisory, let them know the expected timeline. If the reporter had initially set a disclosure deadline, update them if you need more time and negotiate if necessary. Most researchers are willing to extend time if they see progress and good faith from the vendor. Discuss whether the reporter would like to be credited in the advisory or if they prefer anonymity. Also, agree on the date and time of public disclosure. Ideally, release the advisory simultaneously with or shortly after releasing patched software so that users can immediately protect themselves.
- **Post-Resolution Follow-up**: After the fix is released and the advisory is public, it's courteous to follow up with the reporter. Thank them again for their responsible disclosure and perhaps share any lessons learned or improvements you plan because of their report. Keeping a positive relationship can lead to the reporter helping you again in the future or becoming an advocate for your project's security posture.

Throughout all communications with reporters, maintain a professional, collaborative, and appreciative tone. Remember that the broader security community will judge vendors by how they treat researchers – being responsive and fair will enhance your reputation.

- **Communication with Users and Stakeholders:** Users of your software need to know about vulnerabilities in order to protect themselves. However, user communication is typically done **after** or at the time a fix is available to avoid alerting attackers. Key points for user communication include:

- **Security Advisories**: As detailed in the next section, a security advisory is the primary vehicle to inform users. Ensure the advisory is easily accessible – for example, via a dedicated security page on your website, a mailing list announcement, release notes, or a blog post on your official blog. In an enterprise setting, you might also directly email affected customers or issue a press release for very critical issues.

- **Urgency and Guidance**: Clearly communicate how urgent the issue is and what users should do. If it's a critical vulnerability, the advisory should encourage immediate updating. If there are mitigations or workarounds, spell them out so users who can't patch immediately can still reduce risk. Always prefer actionable guidance – for example, "Upgrade to version 4.2.1 or later, which contains the fix" or "Apply the patch linked here" or "As a temporary workaround, disable the XYZ feature until you can update."

- **Clarity and Honesty**: Avoid downplaying the issue or burying the information. Be transparent about what could happen if the vulnerability is exploited but also avoid excessive fear. Use a factual and helpful tone to describe the nature of the vulnerability and its impact in general terms. For example, note, "A buffer overflow in image processing library could allow code execution and allow an attacker to potentially take control of the application". Clarify the affected versions and which versions contain the fix. Explain if only certain configurations are vulnerable so users can assess their own risk.

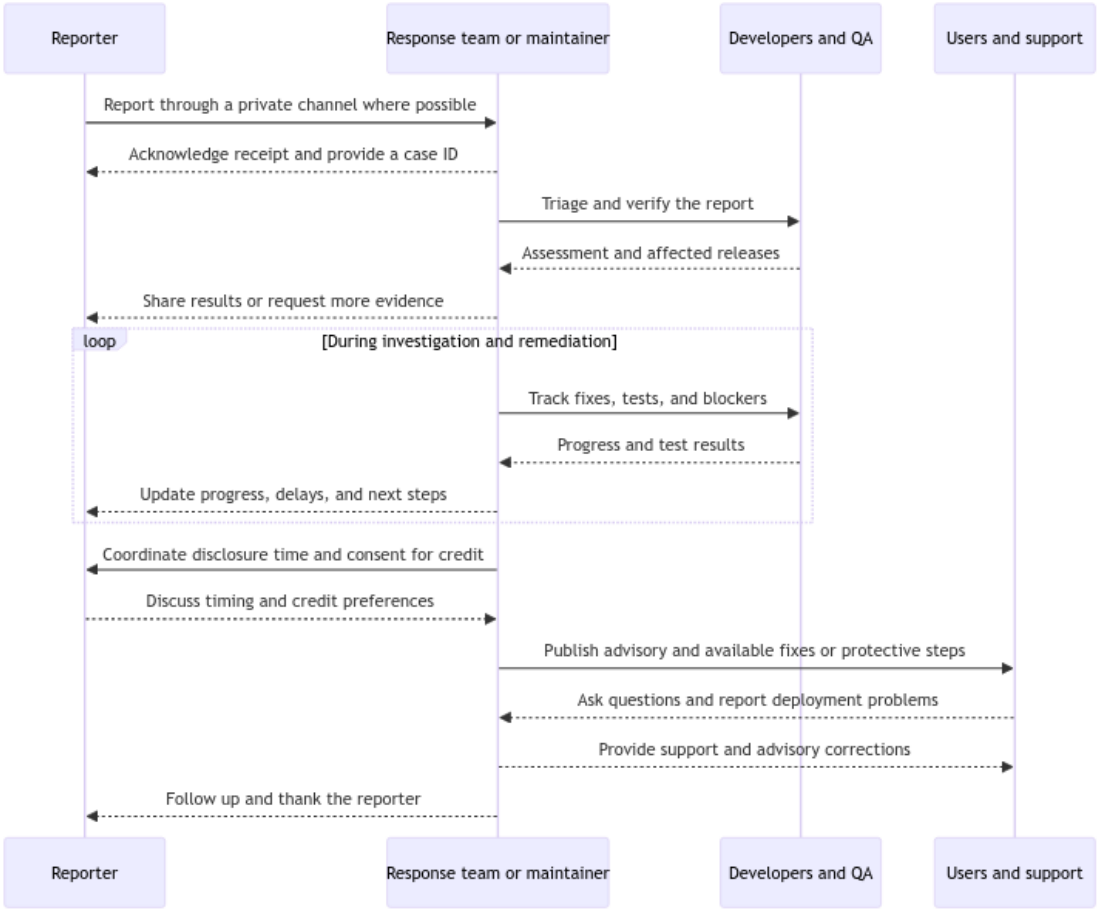
- **Support Channels**: Provide users with a way to get help or ask questions about the issue. This could be your normal support channel or a forum where they can seek guidance if the update process is unclear. Enterprise customers, for example, might have account managers to contact. Open source projects might use their issue tracker or mailing list for follow-up questions. Monitor these channels after disclosure to clarify any confusion.

- **Internal Stakeholders**: Inform internal groups as needed. Customer support teams should be briefed on the issue as soon as or slightly before it's public so they can handle inquiries. Sales or account reps might need advance notice if they are dealing with customers who require prompt notification of security issues. In some cases, such as when the vulnerability must be reported to regulators under certain laws, legal or compliance teams should know details of the vulnerability and its fix. Having a prepared statement or FAQ for internal teams helps ensure a consistent message.

Effective communication promotes a responsive, transparent, and empathetic relationship with both reporters and users. It helps prevent misunderstandings from reporters feeling ignored and users lacking information about security issues. Good communication practices ultimately lead to a smoother coordinated disclosure and a safer environment for everyone.

1.7.1 Communication Overview

This sequence shows the usual communication milestones. Progress updates continue even when there is no new result. An earlier public warning may be needed before a fix is ready; agree reporter credit or anonymity before publication.



1.8 Coordinated Vulnerability Disclosure

Coordinated Vulnerability Disclosure (CVD) is the practice of working together with all involved parties to handle a vulnerability privately until a public disclosure is made at an appropriate time. Embracing CVD principles is highly recommended, as it strikes a balance between **security** (giving vendors time to fix issues) and **transparency** (eventually informing the public). Here's how to approach coordinated disclosure:

- **Private Collaboration First**: When a vulnerability is reported, both the reporter and the vendor agree to address it confidentially before making details public. The vendor commits to investi-

gate and fix the issue while the reporter agrees to withhold public disclosure for a reasonable period or until the fix is released. This cooperation ensures users are not put at undue risk by early disclosure of the bug.

- **Setting a Disclosure Timeline**: An essential aspect of CVD is agreeing on how long the vendor can take to produce a fix before the vulnerability information might be revealed. Many organizations follow an informal 90-day or 60-day guideline, which means that the researcher might publicly disclose the vulnerability anyway if a fix isn't ready after the elapsed time. However, timelines can be flexible if both parties communicate. As a vendor, try to estimate and propose a timeline to issue a patch based on initial risk assessments and complexity of the vulnerability. If you realize you need more time, inform the reporter as soon as possible and provide justification. Researchers often appreciate updates and can grant extensions if progress is evident. The key is to avoid leaving the reporter without guidance; lack of communication is a primary reason researchers go public out of frustration.

- **Involving a Coordinator**: Sometimes involving an impartial third party (coordinator) is useful or necessary. For example, if the reporter can't get a response from the vendor, they might reach out to a CERT or other coordinator to alert the vendor. Conversely, a vendor receiving a report about another vendor's product or a library with a flaw should pass that information to the right party and possibly involve a coordinator for multi-party issues. In multi-vendor situations, such as a vulnerability in an open source component that affects many downstream projects, coordination ensures that everyone gets critical information needed to fix their part and disclosure can happen jointly. Multi-party coordinators can establish a communication channel through mailing list or calls with all vendors and set ground rules for information sharing and the disclosure date. Organizations like CERT/CC or industry groups often help orchestrate this.

  Agree who coordinates the case, which parties need advance notice, what information they may share, and a disclosure date and time with a time zone. Give downstream maintainers enough information to prepare and test their own updates. Revisit the plan together when a dependency delays release or exploitation changes the risk; avoid leaving an embargo open-ended while waiting for every party.

- **No-Fix Scenarios**: Ideally, every reported vulnerability is fixed before disclosure. However, a vendor may decide not to fix a reported issue after deeming it an acceptable risk or deciding they won't fix it due to architectural reasons. In coordinated disclosure, the vendor must communicate this decision to the reporter. They may still choose to disclose it publicly if they believe users should know. As a vendor, you should explain your reasoning in the advisory if this happens. Not fixing is generally discouraged for any significant security problem. Coordinated disclosure in such cases might break down, and the reporter could publish their findings. To maintain goodwill, these situations should be handled with transparency and respect for the

researcher's perspective.

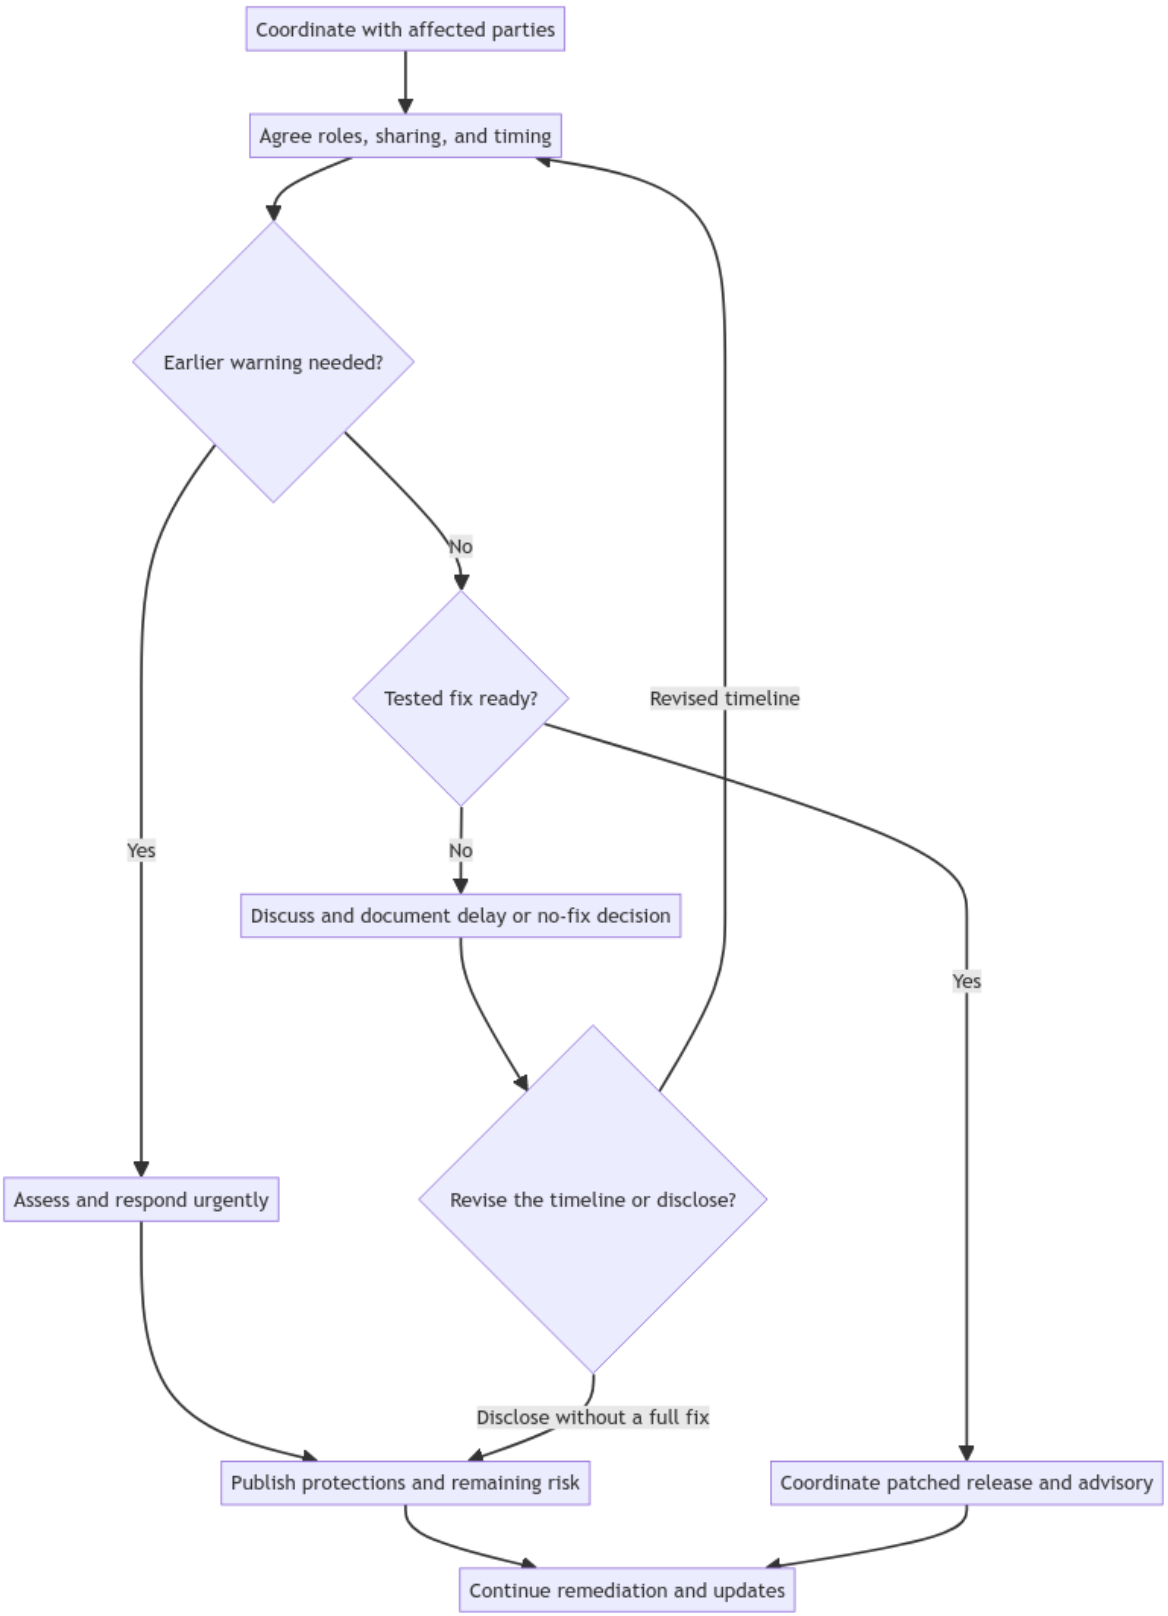
For a deferred fix, record the reason, remaining risk, temporary protections, decision owner, and a date or trigger for review. For an unsupported product, explain the support status and available migration or replacement options. Neither a workaround nor a decision to accept risk should be described as a full fix.

- **Handling Leaks or Early Disclosure**: Despite best efforts, sometimes vulnerability details leak or a reporter publishes early. If vulnerability information becomes public before a fix is out, shift into incident response mode. Quickly assess the risk to users and consider releasing mitigation instructions or interim patches. Communicate openly with users about the situation – even if it's uncomfortable, it's better to acknowledge an early disclosure and provide guidance than to stay silent. Afterwards, analyse what went wrong in the coordination and root causes for early disclosure or information leakage. Then, improve the process. These scenarios underscore why establishing trust and acting swiftly on reports is so important in CVD.

Coordinated vulnerability disclosure depends on **trust and timing**. Vendors should show reporters that they take issues seriously and will act, while reporters give vendors a reasonable opportunity to investigate and respond. Coordination helps users receive a fix or useful protective advice when they learn about a vulnerability, including when a full fix is not yet available.

1.8.1 Coordinated Disclosure Decision Overview

The response team, reporter, and affected parties revisit timing as risk and readiness change. Involve a coordinator where needed for multi-party cases or communication problems. A CVD embargo does not postpone any applicable statutory reporting deadline.


Coordinate with affected parties
Agree roles, sharing, and timing
Earlier warning needed?
No
Tested fix ready?
Revised timeline
Yes
No
Discuss and document delay or no-fix decision
Yes
Assess and respond urgently
Revise the timeline or disclose?
Disclose without a full fix
Publish protections and remaining risk
Coordinate patched release and advisory
Continue remediation and updates

1.9 Publishing Advisories

Publishing a security advisory is the capstone of the vulnerability handling process. An advisory should explain what the issue is, who is affected, and how it can be fixed. These answers support a public record of the vulnerability and its fix, and it informs all users about what action to take. A well-crafted advisory considers the following in its timing and content.

- **Timing of Advisory Release**: Coordinate the advisory release with the availability of the fix. Ideally, the advisory should be published at the same time (or very shortly after) you release the patched software version. This way, users reading the advisory can immediately take action to secure their systems. If exploitation, a leak, or the absence of a fix makes earlier warning necessary, publish the known risk and protective steps while clearly stating what remains unresolved. Avoid delaying an advisory long after the fix, as users might not realize a security update is important without the context. If detailed technical information is delayed to allow deployment, record why and tell users when to expect it. In cases where a fix is rolled out silently (e.g. via auto-update), an advisory should still follow to document the issue.
- **Summary**: Provide a brief description of the vulnerability and its impact. For example, "A buffer overflow in the image parsing library could allow an attacker to execute arbitrary code on versions 1.2 through 1.4 of the application."
- **Affected Products/Versions**: List the specific product names and versions that are vulnerable. Be as precise as possible (e.g. "Versions 1.0.0 to 1.4.2 are affected; version 1.4.3 and above contain the fix"). If older, unsupported versions are also affected, mention them as well but note if they are not patched.
- **Solution (Fixed Versions)**: State the version numbers or update that fixes the issue. If the fix is available as a patch, hot-fix, or commit, provide links or instructions on how to obtain and apply it. In open source projects, this might be a reference to a specific commit or a new release tag.
- **Third-Party Components**: Link the upstream advisory or fix where available, and explain what it means for your product: affected releases, the release containing the updated component, and any action users still need to take. A link to an upstream fix alone may not tell users whether your shipped product has been corrected.
- **Publication and Updates**: Include the initial publication date, the date of the latest revision, a short description of important changes, and a contact for questions or corrections. Keep a stable advisory URL so users and automated tools can find later updates.
- **Workarounds/Mitigations**: If applicable, describe any temporary measures users can take if they cannot immediately apply the fix. For example: "Until you can update, you can disable the image parsing feature by doing X," or "Configure the firewall to block port Y to mitigate the

issue." Not all advisories have mitigations, but include them if they exist.

- **Acknowledgement**: Credit the reporter or others involved in discovering the vulnerability if they consent to be named. This typically goes near the end: e.g. "We thank Jane Doe for reporting this issue responsibly."
- **CVE or GCVE or any Tracking ID**: If a CVE or GCVE ID has been assigned to this vulnerability, include it. CVEs or GCVE are useful for indexing the issue in global databases and for users who track vulnerabilities via scanning tools. If you don't have a CVE or GCVE ID, you should include an internal tracking number instead of nothing at all – CVEs or GCVE are not strictly required but are standard for notable vulnerabilities. For more details, GCVE-BCP-04 - Recommendations and Best Practices for ID Allocation.
- **Advisory Content** – Additional Information: In addition to the basics, including more details can be helpful:
  - **Severity/Rating**: Indicate how severe the issue is (Critical High/Medium/Low) to help users prioritize the update. If you use CVSS, provide the score and vector string.
  - **Discovery Timeline**: It is a best practice to include a timeline of events and describe when the issue was reported, when it was fixed, and when the advisory is published. This transparency can demonstrate your adherence to a responsible process and give credit to the timeline of the reporter's cooperation.
  - **Technical Details**: Depending on your audience, you might add a section with more technical explanation of the vulnerability. This can help advanced users or peer reviewers understand the nature of the flaw. It's also useful for historical record and for other developers to learn. If the vulnerability is complex, you might summarize how it was found or how it works. You may omit deep technical details if you fear it might aid exploit development. However, sharing details of a vulnerability after a patch is issued is considered good for transparency and education.
  - **References**: Link to any relevant references. For example, if this issue was discussed publicly or if it's related to a known vulnerability class, you might reference external documents or prior advisories.
  - **Update Instructions**: If updating is non-trivial and requires configuration changes or a series of steps, outline the steps or link to upgrade documentation.
  - **Format and Distribution**: Present the advisory in a format accessible to your users. Many organizations use plain text or Markdown for advisory documents and sometimes publish advisories in HTML on websites. Use a consistent template so users know where to find information. Consider distributing the advisory through multiple channels:

    * Post it on your official website (preferably in a dedicated Security Advisories section). A machine-readable format should be available to facilitate the discovery and processing of vulnerabilities. This can be achieved with open source tool such as vulnerability-lookup. For more details about machine-readable format, GCVE-BCP-05 - GCVE Vulnerability Format (Updated CVE Record Format).
    * Machine-readable security advisories should be made available under an open distribution license—such as an Open Source–approved license or a permissive Creative Commons license (e.g., CC-BY or CC0)—to ensure efficient redistribution, reuse, and automated processing across ecosystems.
    * Send it to a mailing list dedicated to announcements or security updates.
    * Publish it via your project's blog or news section.
    * Share on forums or community channels where users gather.
    * For open source projects, you might also use your source repository's advisory features (e.g. GitHub Security Advisory) which can send alerts to users of the project.

  Ensure that once published, the advisory remains available indefinitely for reference. Don't take down old advisories; they serve as a historical security record.

- **Post-Publication Monitoring**: After publishing, monitor the reaction. Be ready to answer questions from users or the media. If any detail in the advisory is found to be incorrect or unclear, issue a correction or update the advisory. Sometimes after release, additional research may discover that the vulnerability affected additional versions or that the fix had a bug. These discoveries and their subsequent fixes require an update to the advisory. Monitor for any signs of public exploitation now that the vulnerability is published. If something arises, you might need to alert users or provide additional mitigation advice.

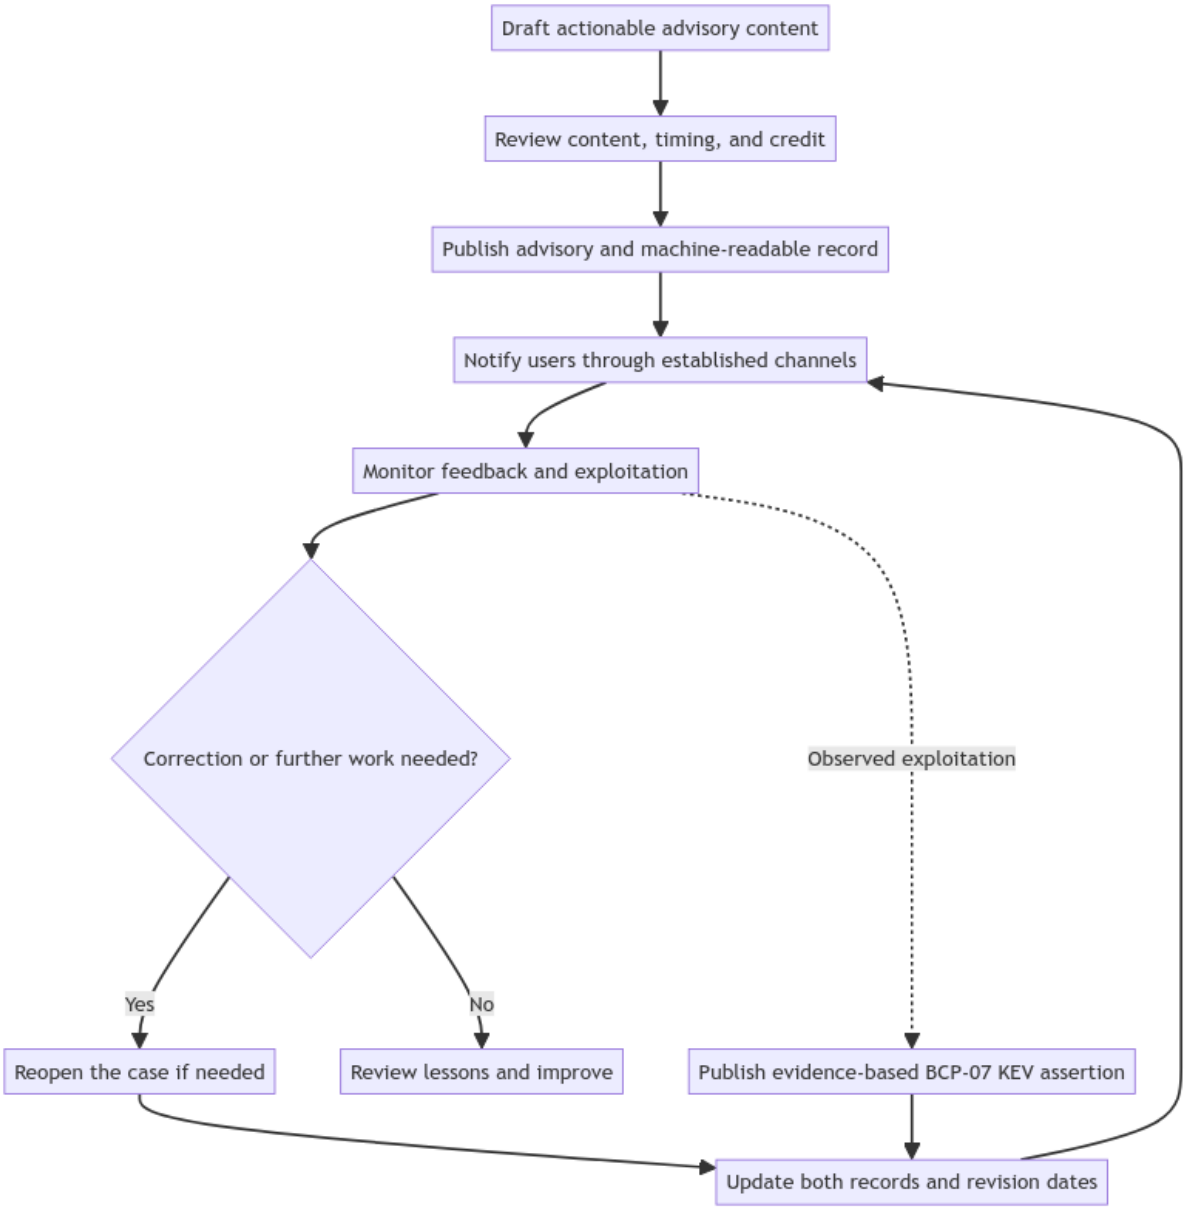
  - Distinguish observed exploitation from an assessment that exploitation is possible. Document the evidence and uncertainty in the advisory. When asserting known exploitation, publish a corresponding Known Exploited Vulnerability (KEV) assertion using the GCVE BCP-07 format to enable consistent machine-readable consumption.
  - Check that users can obtain and apply the update, and use available support feedback or deployment results to detect failures, regressions, or incomplete fixes. Respect user privacy when collecting such information. Reopen the case when necessary, repeat the affected verification and fix steps, and update both the advisory and machine-readable record. Releasing a package does not by itself confirm that deployed systems are protected.

Publishing advisories is a fundamental duty to your user base. A well-handled advisory not only helps users protect themselves but also builds your reputation for transparency. Stakeholders from enterprise customers to open source users will appreciate clear and timely advisories. Remember, acknowl-

edging security issues publicly does not mean your software is less secure. On the contrary, it shows you take security seriously and deal with issues head-on, which ultimately **increases** user trust.

1.9.1 Advisory Publication and Follow-up Overview

Publish consistent human-readable and machine-readable information, then keep both current. Follow-up checks whether users can apply the protection and whether new evidence changes the assessment.



1.10 Best Practices

The following best practices reinforce the guidance above and provide a quick checklist for building a robust process:

- **Encourage Responsible Reporting**: Make it easy for people to report vulnerabilities to you. Publish clear instructions and promise a supportive response. Consider using a security.txt file on your website or repository so automated tools and researchers can find your contact info. Provide a way to encrypt sensitive reports (PGP key). Prominently assure researchers that you welcome reports and will not take legal action for good-faith efforts.
- **Respond and Remediate Promptly**: Time is of the essence in security. Strive to acknowledge reports quickly and fix critical issues as fast as possible without sacrificing quality. Even if a bug is complex, make interim plans and keep all parties informed. A slow response can result in public disclosure without a fix, which puts users at risk. Treat researchers as partners in security. Be polite, thankful, and honest in all interactions.
- **Protect Sensitive Information**: Handle vulnerability details on a need-to-know basis until disclosure. Use encrypted channels for communications. Limit public discussion of the issue until the advisory is released. If working with multiple organizations, use Non-Disclosure Agreements (NDA's) or trusted channels as needed to prevent leaks. Always error on the side of caution with sensitive information. Use a clear classification standard, while distributing the information or interacting with third-parties, such as the Traffic Light Protocol (TLP).
- **Review AI-Assisted Outputs**: If AI-assisted tools are used to classify, enrich, summarize, prioritize, or draft vulnerability information, their outputs should be considered provisional until reviewed by a qualified human. AI-assisted severity ratings, affected-version claims, exploitability assessments, KEV assertions, remediation guidance, and advisory text must not be published as authoritative without human validation. Where possible, machine-readable advisories should preserve provenance for AI-assisted fields, including the tool/model used, confidence score, timestamp, and review status. Non-public vulnerability information must not be submitted to external AI systems unless confidentiality, retention, training-use, and legal requirements have been assessed.
- **Coordinate and Collaborate**: Follow coordinated disclosure principles. Work with reporters on timelines. If multiple vendors are involved, take initiative in reaching out to coordinate a joint response. Share information with platform owners or other stakeholders if the vulnerability could extend to them. Security is a team effort; collaboration can be the difference between a minor contained issue and a major public incident.
- **Learn and Improve**: After each vulnerability case review process steps that went well or could be improved. Implement improvements to your systems and practices based on these obser-

vations. Each vulnerability is also a lesson for developers: share the root cause and teach the team how to avoid that class of bug in the future. Over time, your software should get more secure and your handling process more efficient.

Review practical measures such as time to acknowledge, time to release a fix, overdue cases, recurring weaknesses, and failed or incomplete updates. Use the results to improve tests, ownership, and update delivery. Review the public policy alongside the internal workflow so promised response times and actual practice stay consistent.

- **Stay Informed on Security Practices**: The security landscape evolves. Keep your vulnerability handling aligned with current best practices including industry standards and new guidelines from organizations like CERT/CC, FIRST.org, or NIST. Participate in security communities or forums to learn from others' disclosure experiences. For open source maintainers, many organizations and groups provide resources and forums to discuss disclosure challenges unique to open source – leveraging these can strengthen your own procedures.
- **Legal and Regulatory Compliance**: Be aware of any legal requirements regarding vulnerability disclosure that apply to you. Certain industries or regions may have laws about reporting breaches or vulnerabilities to regulators. While this guide focuses on the process itself, always ensure your legal counsel is in the loop if a vulnerability could trigger regulatory obligations. Also, having a clear policy with safe harbour language can protect both you and researchers by setting mutual expectations.
- **Recognize and Reward**: When possible, acknowledge the contributions of those who help improve your security. Publicly credit researchers (with their permission) in advisories or on a Hall of Fame page to encourage others to report. If resources allow, consider offering rewards.

By adhering to these best practices, organizations and open source projects can create a robust ecosystem for vulnerability handling and disclosure. The result is a win-win: security researchers have confidence that reporting issues will lead to constructive action, and users benefit from timely fixes and open communication about the security of the products they rely on.

By following the guidance in this document, software maintainers and organizations can better protect their users and improve their products' security over time. Vulnerability handling and disclosure is an ongoing commitment – one that, when done right, significantly reduces risk and builds trust in the software. It transforms security vulnerabilities from potential disasters into opportunities for learning and strengthening the resilience of our systems.

2 Appendix A: EU Cyber Resilience Act Reporting

This appendix provides complementary operational context for organizations subject to the European Union Cyber Resilience Act (CRA), Regulation (EU) 2024/2847. It does not change the general CVD workflow described in this BCP and is not legal advice. Organizations should assess the CRA's scope, definitions, exceptions, and application to their products and roles with qualified legal counsel.

2.1 What the CRA Adds

The CRA establishes horizontal cybersecurity requirements for products with digital elements made available on the EU market. Among other lifecycle obligations, it requires manufacturers to handle vulnerabilities effectively. This includes documenting relevant aspects, providing security updates during the support period, maintaining a coordinated vulnerability disclosure policy, facilitating the sharing of vulnerability information, and meeting the mandatory reporting obligations in Article 14.

The Article 14 reporting obligations apply from **11 September 2026**, ahead of the CRA's general application from 11 December 2027. A manufacturer must report when it becomes aware of either:

- an **actively exploited vulnerability** contained in its product with digital elements; or
- a **severe incident having an impact on the security** of its product with digital elements.

These are legal notification duties, not optional vulnerability-publication practices. They apply to the manufacturer when the applicable CRA threshold is met. A vulnerability does not become reportable merely because it has been reported by a researcher or assigned a vulnerability identifier: the manufacturer must assess the facts against the CRA definitions and criteria. Conversely, assigning a CVE or GCVE identifier, publishing an advisory, notifying a coordinator, or following the CVD steps in this BCP does not by itself satisfy Article 14.

2.2 Relationship to Coordinated Vulnerability Disclosure

CRA reporting and CVD are **complementary processes with different purposes and recipients**:

- **CVD** coordinates the private receipt, validation, remediation, and appropriately timed public disclosure of vulnerability information among a reporter, manufacturer or maintainer, coordinators, and other affected parties.
- **CRA reporting** is a mandatory regulatory notification by the manufacturer to the designated authorities through the Single Reporting Platform (SRP) when an Article 14 reporting threshold is met.

The regulatory deadlines should therefore be tracked in parallel with the normal CVD timeline. A manufacturer should continue to protect embargoed information, work with the reporter and affected downstream parties, prepare and distribute fixes, and publish an advisory at the coordinated time. Submitting an SRP notification is not a public advisory, does not remediate the vulnerability, and does not replace communication with the reporter or users. Similarly, an agreed CVD embargo cannot postpone a statutory CRA deadline.

Researchers and other external reporters should continue to use the manufacturer's published security contact or CVD channel. Article 14 places this reporting obligation on manufacturers; the SRP is not a replacement intake channel through which researchers report vulnerabilities to manufacturers. A reporter can help the manufacturer make a timely assessment by clearly indicating any evidence of malicious exploitation, observed incidents, affected products and versions, and the time at which that evidence was observed.

2.3 The Single Reporting Platform

The CRA requires ENISA to establish and maintain the SRP as the electronic reporting mechanism for Article 14 notifications. The platform is designed so the manufacturer can notify the CSIRT designated as coordinator for CRA purposes and ENISA through one submission. The manufacturer's selected notifying endpoint determines the initially notified CSIRT; the notification is then made available simultaneously to ENISA. The CRA provides rules for selecting that endpoint, principally based on the manufacturer's main establishment in the EU and, where that does not resolve the choice, other connecting factors specified in Article 14.

The SRP is intended to support confidential, time-sensitive regulatory reporting. It is distinct from public vulnerability databases, advisory feeds, and identifier-allocation systems. Manufacturers should consult ENISA's Single Reporting Platform frequently asked questions for current practical guidance about access, account preparation, submission, roles, and operation of the platform. Because the platform and the authorities' implementation guidance may evolve, organizations should verify the current instructions rather than encode interface-specific assumptions into their CVD policy.

2.3.1 Reporting stages and deadlines

The deadlines run from the manufacturer's awareness, not from validation completion, patch availability, identifier assignment, or public disclosure. Article 14 establishes staged reporting:

Reportable event	Initial notification	Follow-up notification	Final report
Actively exploited vulnerability	Early warning without undue delay and, in any event, within **24 hours** of awareness	Vulnerability notification without undue delay and, in any event, within **72 hours** of awareness	No later than **14 days after a corrective or mitigating measure is available**
Severe incident affecting product security	Early warning without undue delay and, in any event, within **24 hours** of awareness	Incident notification without undue delay and, in any event, within **72 hours** of awareness	No later than **one month after the incident notification**

The staged model permits the manufacturer to provide an early warning while facts are still developing and then supplement it with a fuller assessment. The required content differs by stage and event. For an actively exploited vulnerability, for example, the later reporting includes available information about the vulnerability, its exploitation, affected products, and corrective or mitigating measures. Manufacturers should use the current SRP workflow and ENISA guidance to determine the fields required for each submission.

Article 14 contains limited rules concerning when information may be omitted or submitted later, including where operational security reasons justify delaying certain details. These provisions should not be treated as a general extension of the deadlines. Document the justification and seek appropriate legal or authority guidance when relying on an exception.

2.4 Integrating CRA Reporting into a CVD Process

Manufacturers in or serving the EU market should add a CRA decision path to their existing vulnerability-handling procedure without weakening that procedure:

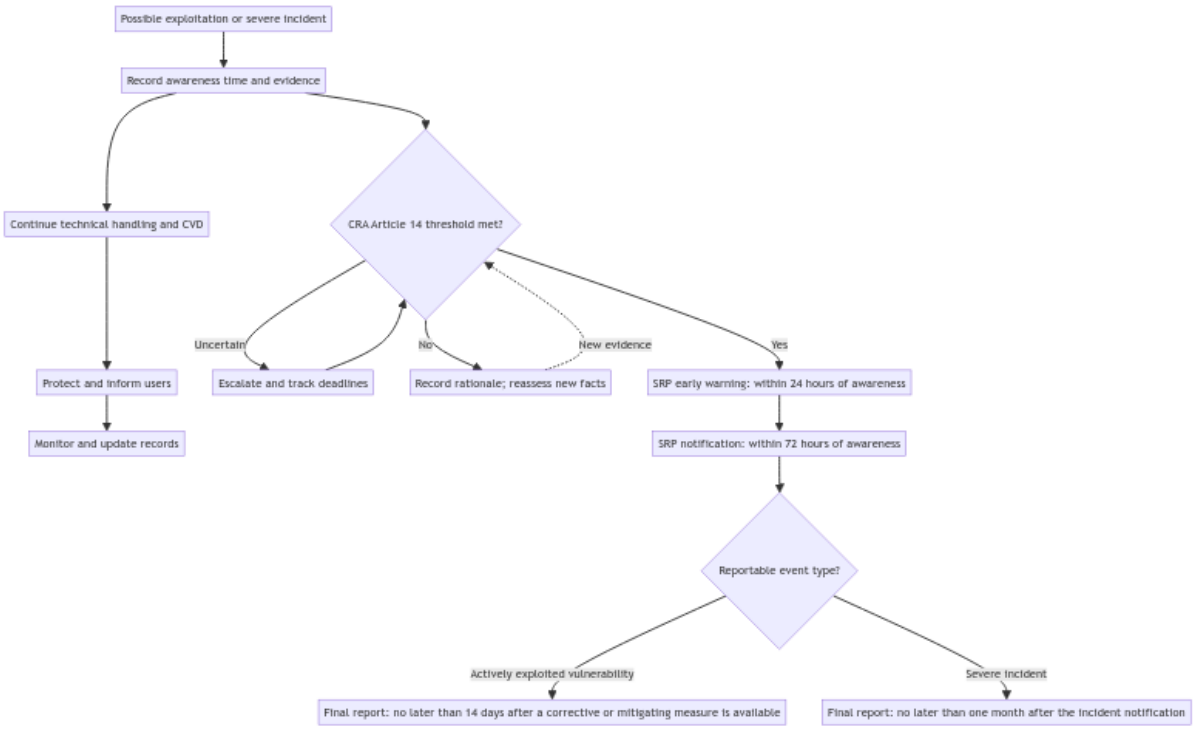
1. **Assign ownership before a report arrives.** Identify the legal manufacturer, the team authorized to submit through the SRP, its backups, and the applicable notifying endpoint. Prepare the necessary platform access and keep contact information current.

2. **Record awareness time.** When information might indicate active exploitation or a severe product-security incident, record when the manufacturer first became aware and preserve the supporting evidence. Do not wait for a complete root-cause analysis before starting the reporting assessment.
3. **Triage both tracks.** Continue technical validation and CVD coordination while an authorized legal/product-security function promptly determines whether the CRA threshold is met. Escalate uncertainty early because the first deadline may be 24 hours.
4. **Submit and update through the SRP.** Meet each applicable stage and retain submission receipts, case references, submitted content, timestamps, decisions, and the evidence supporting them. Protect access to the report and observe applicable confidentiality rules.
5. **Continue remediation and CVD.** Develop and test the correction or mitigation, coordinate with the reporter and affected parties, request an identifier where appropriate, notify users, and publish a machine-readable advisory. Do not regard the regulatory notification as closure of the vulnerability case.
6. **Link records without conflating them.** Internally associate the SRP case, vulnerability identifier, advisory, affected-product inventory, remediation work, and CVD correspondence. Keep each system's audience, disclosure status, and access controls explicit.
7. **Exercise the workflow.** Test the 24-hour escalation path, including weekends and absences, and periodically review it against the current CRA text, national CSIRT guidance, and ENISA's SRP FAQ.

A practical case record can include fields such as `CRA assessment required`, `awareness time`, `active-exploitation evidence`, `severe-incident assessment`, `reporting decision and rationale`, `SRP endpoint`, `24-hour/72-hour/final-report deadlines`, `submission references`, and `corrective or mitigating measure availability time`. These fields supplement—not replace—the CVD tracking data recommended elsewhere in this BCP.

2.4.1 Parallel CRA Reporting and CVD Overview

For manufacturers subject to the CRA, record awareness time and assess Article 14 promptly while technical handling continues. The reporting stages below summarize the preceding table: the 24-hour and 72-hour limits both run from awareness and require reporting without undue delay. The two final-report deadlines have different triggers. Retain submission receipts and link the regulatory and CVD records with appropriate access controls.


Possible exploitation or severe incident
Record awareness time and evidence
Continue technical handling and CVD
CRA Article 14 threshold met?
Uncertain
No
New evidence
Yes
Protect and inform users
Escalate and track deadlines
Record rationale; reassess new facts
SRP early warning: within 24 hours of awareness
Monitor and update records
SRP notification: within 72 hours of awareness
Reportable event type?
Actively exploited vulnerability
Severe incident
Final report: no later than 14 days after a corrective or mitigating measure is available
Final report: no later than one month after the incident notification

3 References

- European Union - Regulation (EU) 2024/2847 (Cyber Resilience Act)
- ENISA - Single Reporting Platform: Frequently Asked Questions
- OWASP - Vulnerability Disclosure - OWASP Cheat Sheet
- CIRCL - Coordinated Vulnerability Disclosure (CVD) Policy
- NIST - Recommendations for Federal Vulnerability Disclosure Guidelines
- Security Engineering — Third Edition - Ross Anderson
- IETF Internet-Draft - Responsible Vulnerability Disclosure Process draft-christey-wysopal-vuln-disclosure-00

4 Acknowledgements

4.1 BCP-02 Coordinators

- Alexandre Dulaunoy, CIRCL
- Sascha Rommelfangent, CIRCL

4.2 Contributions

The GCVE initiative gratefully acknowledges the substantial contributions from the following individuals:

- Cédric Bonhomme
- Xavier Claude
- Matthew J. Harmon
- Quentin Jerome
- Tony Rutkowski
- CSIRTs and PSIRTs Community